

# Inverse Pricing for the Implied Volatility Surface using Physics-Informed Neural Networks

Alex Yang

December 6, 2025

## Abstract

This project studies the inverse problem of recovering a smooth, continuous and arbitrage-free implied volatility (IV) surface  $\sigma(K, T)$  from sparse and noisy market prices of European options. I propose a Physics-Informed Neural Network (PINN) framework in which a neural network parameterizes the IV surface, the Black–Scholes formula serves as a differentiable forward map from volatility to price, and no-arbitrage constraints are embedded as soft penalties in the loss. Using SPY options data from OptionMetrics (WRDS), I show that a two-phase “warm-up” training strategy is essential: first fit prices without constraints, then gradually introduce calendar- and butterfly-arbitrage penalties. The resulting surface achieves high pricing accuracy (RMSE around \$1.76 on the calibration day), zero calendar-arbitrage violations on a dense grid, and strong zero-shot and fine-tuned generalization to neighboring trading days.

## 1 Introduction and Background

The implied volatility (IV) surface  $\sigma(K, T)$  is a central object in equity and index options markets. For each strike  $K$  and time-to-maturity  $T$ , the Black–Scholes formula maps a volatility input to a theoretical option price. In practice, the market quotes only a finite and noisy set of mid prices  $C^{\text{mkt}}(K_i, T_i)$ ; the IV surface itself is not directly observable. Recovering a smooth, economically meaningful  $\sigma(K, T)$  from these discrete prices is therefore a classical ill-posed inverse problem.

A naive approach inverts the Black–Scholes formula pointwise to obtain a cloud of implied volatilities  $\hat{\sigma}_i$  and then interpolates across  $(K, T)$ . This is unstable in regions with few quotes, particularly for short maturities, deep in-the-money (ITM), or deep out-of-the-money (OTM) options, and typically produces jagged surfaces that generate spurious arbitrage. From an inverse-problems perspective, the map

$$\sigma \mapsto \{C^{\text{BS}}(K_i, T_i; \sigma)\}_i$$

is compact and poorly conditioned; regularization is not optional but essential.

In this project, I frame IV-surface calibration as a regularized inverse problem solved via a Physics-Informed Neural Network (PINN). A multilayer perceptron (MLP) parameterizes  $\sigma_\theta(K, T)$ ; a differentiable Black–Scholes layer turns this into model prices  $C_\theta(K, T)$ ; and the training loss balances data fidelity against financial “physics” in the form of no-arbitrage constraints:

$$\min_{\theta} \mathcal{L}(\theta) = L_{\text{price}}(\theta) + \lambda_{\text{cal}} L_{\text{calendar}}(\theta) + \lambda_{\text{but}} L_{\text{butterfly}}(\theta). \quad (1)$$

The calendar term penalizes violations of the requirement that total variance  $\sigma^2 T$  be non-decreasing in  $T$ ; the butterfly term penalizes violations of convexity of the call price in  $K$ , which is equivalent to non-negativity of the risk-neutral density.

Empirically, I find that directly enforcing these constraints from epoch 0 often causes the network to collapse to an almost flat surface that trivially satisfies derivative-based penalties while ignoring the data. A key methodological contribution of this project is a simple but effective “warm-up” strategy: train first on prices only, then smoothly turn on the physics penalties with small weights.

## Contributions

The main contributions of this project are:

- I implement an end-to-end PINN framework for IV-surface inversion where the network outputs implied volatility (rather than prices) and arbitrage constraints are evaluated through the Black–Scholes map.
- I design and evaluate calendar- and butterfly-arbitrage penalties based on total variance monotonicity and call-price convexity.
- I propose and validate a two-phase warm-up training schedule that prevents mode collapse and yields stable convergence.
- On SPY options from OptionMetrics, I demonstrate that the trained model is arbitrage-free on a dense grid, achieves competitive pricing accuracy on the calibration day, and generalizes well to neighboring trading days, with efficient transfer learning through short fine-tuning runs.

## 2 Literature Review

### 2.1 Nonparametric IV-surface calibration

Early work views volatility-surface recovery as a nonparametric inverse problem. A seminal example is Aït-Sahalia and Lo (1998), who study the nonparametric estimation of state-price densities and implied-volatility surfaces from cross-sections of option prices. They emphasize that the mapping from prices to implied quantities is ill-posed and highly sensitive to noise. To regularize the problem, they minimize a data misfit plus a smoothness penalty:

$$\min_{\sigma} \sum_i (C^{\text{BS}}(K_i, T_i; \sigma) - C_i^{\text{mkt}})^2 + \lambda \mathcal{R}(\sigma),$$

where  $\mathcal{R}(\sigma)$  penalizes curvature in  $K$  and  $T$ . The regularization parameter  $\lambda$  trades off bias and variance. Empirical results show that this smoothing yields visually plausible surfaces and more stable Greeks than pointwise inversion. Conceptually, this work clarifies that regularization is the right lens for IV-surface estimation; in my project, the PINN penalties play a similar role to  $\mathcal{R}(\sigma)$  but are tailored to financial no-arbitrage.

A complementary line of research uses constrained smoothing splines to impose no-arbitrage conditions directly on option prices or implied volatilities. For instance, Laurini (2010) construct smoothed implied-volatility curves that satisfy monotonicity and convexity constraints using constrained spline fitting. These methods show that shape constraints can be enforced efficiently at the curve level, but extending them to the full two-dimensional  $(K, T)$  surface while preserving tractability is non-trivial. My approach inherits the spirit of constraint-based smoothing but leverages automatic differentiation and neural networks to scale to a full surface.

## 2.2 Parametric volatility surfaces and local volatility

In practice, parametric models such as SVI (Stochastic Volatility Inspired) are popular because they encode no-arbitrage conditions in a finite set of parameters. Gatheral and Jacquier (2014) analyze arbitrage-free SVI parameterizations and derive conditions under which an SVI smile is free of static arbitrage in strike and time. While these models are robust and interpretable, their global parametric structure can limit flexibility when the market surface exhibits local irregularities. My neural-network parameterization can be viewed as a highly flexible, data-driven alternative; the no-arbitrage penalties in my loss function play a similar role to the analytical SVI constraints.

Local volatility models, dating back to Dupire (1994), provide another route: the local-volatility surface  $\sigma_{\text{loc}}(S, t)$  can be recovered from option prices via the Dupire equation. However, calibration is extremely ill-posed and sensitive to noise and interpolation choices; see, for example, discussions in Cont and Tankov (2004). Although I do not work in the local-volatility framework, these works reinforce the view that inverse option problems must be treated with appropriate regularization and constraints.

## 2.3 Deep learning and PINNs for option pricing

More recently, neural networks have been proposed both as fast approximators for pricing functions and as flexible parameterizations of risk-neutral dynamics. For example, Hernandez (2017) use feedforward networks to approximate option prices given model parameters, showing significant speedups over Monte Carlo or finite differences once trained. However, such networks typically ignore no-arbitrage constraints and may generate economically invalid predictions when used outside the training domain.

Physics-Informed Neural Networks (PINNs) were introduced by Raissi et al. (2019) as a deep learning framework for solving forward and inverse problems governed by partial differential equations (PDEs). The key idea is to augment the data loss with a PDE residual evaluated at collocation points, encouraging the network’s output to satisfy the underlying physical law. PINNs have since been applied to Black–Scholes-type PDEs and option pricing; for instance, Zhang et al. (2023) train networks to approximate option prices while enforcing the Black–Scholes PDE and payoff boundary conditions.

My work is inspired by these PINN ideas but differs in emphasis. I apply the physics constraints not to a price network solving the PDE, but to an IV-surface network whose outputs are fed through the Black–Scholes formula. The “physics” in my loss are no-arbitrage properties—calendar monotonicity and strike convexity—which are intimately related to PDE structure but easier to evaluate on market-relevant quantities.

## 2.4 Connection to this project

Overall, the literature suggests three guiding principles:

1. IV-surface calibration is an ill-posed inverse problem and must be regularized.
2. No-arbitrage conditions (monotonicity, convexity, positivity of densities) are crucial to obtain economically meaningful surfaces.
3. Neural networks and PINNs provide flexible function approximators with built-in automatic differentiation, well-suited for enforcing such constraints through loss terms.

My project combines these ideas by using a neural-network parameterization, a Black–Scholes forward map, and differentiable no-arbitrage penalties, together with a training schedule inspired by regularization theory: first fit the data, then gradually strengthen the constraints.

## 3 Data and Preprocessing

### 3.1 Data source and selection

I use option data for the SPDR S&P 500 ETF (ticker SPY) from the OptionMetrics database via WRDS. The primary calibration date is 3 January 2023; I also extract data for 4 and 5 January 2023 to evaluate temporal generalization.

For each date, I obtain:

- European call options across a wide range of strikes  $K$  and maturities  $T$ ,
- bid and ask quotes, from which I compute mid prices,
- underlying SPY close price  $S$ ,
- zero-coupon yield curves, from which I interpolate the relevant risk-free rate  $r(T)$ ,
- an estimate of the continuous dividend yield  $q(T)$  for SPY.

### 3.2 Filtering and cleaning

Raw option chains contain many illiquid or economically uninformative contracts. I therefore apply the following filters:

- **Price filter:** discard options with mid price below \$0.01.
- **Maturity filter:** keep options with time-to-maturity between 7 days and 2 years. Very short-dated options are highly sensitive to microstructure noise; very long-dated options are often illiquid.
- **Moneyness filter:** restrict to contracts with moneyness  $K/S$  between 0.7 and 1.3. Deep ITM/OTM options tend to be noisy and exert disproportionate leverage on the calibration.

For each option, I compute the mid price as the simple average of bid and ask:

$$C^{\text{mkt}} = \frac{\text{bid} + \text{ask}}{2}.$$

### 3.3 Feature engineering

For each remaining contract, I construct the following features:

- **Moneyness:**  $m = K/S$ ,
- **Log-moneyness:**  $x = \log(K/S)$ ,
- **Time-to-maturity:**  $T$  in years, computed from calendar dates,
- **Risk-free rate:**  $r = r(T)$  from the interpolated yield curve,
- **Dividend yield:**  $q = q(T)$ .

I standardize inputs  $(x, T)$  to zero mean and unit variance before feeding them into the neural network; this helps optimization.

### 3.4 Reference implied volatility

For diagnostics and visualization, I compute a reference implied volatility  $\hat{\sigma}_i$  for each contract by numerically inverting the Black–Scholes call price:

$$C^{\text{BS}}(S, K, T, \hat{\sigma}_i, r, q) = C_i^{\text{mkt}}.$$

I use a robust root-finding method (Brent’s method) initialized with a reasonable volatility bracket. These  $\hat{\sigma}_i$  are *not* used as training targets, but provide a useful benchmark scatter plot and starting point for understanding the structure of the market IV surface.

## 4 Methodology

### 4.1 Neural network parameterization of the IV surface

I parameterize the implied volatility surface as a neural network

$$\sigma_\theta : (x, T) \mapsto \sigma_\theta(x, T) > 0,$$

where  $x = \log(K/S)$  is log-moneyness. The positive range is enforced by using a Softplus activation at the output.

The architecture, implemented as `IVSurfaceMLP`, is:

- **Input layer:** 2 units (log-moneyness  $x$  and time-to-maturity  $T$ ), after standardization.
- **Hidden layers:** three fully connected layers with widths [128, 128, 64] and SiLU activations.
- **Output layer:** 1 unit with Softplus activation to enforce  $\sigma_\theta(x, T) > 0$ .
- **Dropout:** rate 0.05 on hidden layers for mild regularization.

This MLP is flexible enough to approximate realistic smile and term-structure shapes, while remaining small enough to train efficiently.

### 4.2 Black–Scholes forward map

Given the network output  $\sigma_\theta(x, T)$ , I compute model call prices via the Black–Scholes formula with continuous dividend yield:

$$C_\theta(S, K, T, r, q) = C^{\text{BS}}(S, K, T, \sigma_\theta(x, T), r, q) \tag{2}$$

$$= Se^{-qT} \Phi(d_1) - Ke^{-rT} \Phi(d_2), \tag{3}$$

where

$$d_1 = \frac{\log(S/K) + (r - q + \frac{1}{2}\sigma_\theta^2)T}{\sigma_\theta \sqrt{T}}, \quad d_2 = d_1 - \sigma_\theta \sqrt{T},$$

and  $\Phi$  denotes the standard normal cumulative distribution function. This mapping is implemented in a differentiable way using automatic differentiation.

### 4.3 Physics-informed loss: data and no-arbitrage terms

The total loss function (1) consists of three terms.

### 4.3.1 Price data fidelity

The primary objective is to fit observed mid prices. Over the training set  $\mathcal{D}$  of contracts indexed by  $i$ , I use a mean squared error:

$$L_{\text{price}}(\theta) = \frac{1}{|\mathcal{D}|} \sum_{i \in \mathcal{D}} \left( C_\theta(S_i, K_i, T_i, r_i, q_i) - C_i^{\text{mkt}} \right)^2. \quad (4)$$

### 4.3.2 Calendar-arbitrage penalty

In an arbitrage-free setting, total variance  $w(K, T) = \sigma^2(K, T)T$  should be non-decreasing in  $T$  for each strike:  $\partial_T w(K, T) \geq 0$ . To encourage this, I define a penalty over a grid of collocation points  $(K_j, T_j)$ :

$$L_{\text{calendar}}(\theta) = \frac{1}{|\mathcal{G}_{\text{cal}}|} \sum_{(K, T) \in \mathcal{G}_{\text{cal}}} \left( \min \{0, \partial_T w_\theta(K, T)\} \right)^2, \quad (5)$$

where  $w_\theta(K, T) = \sigma_\theta^2(x, T)T$  and  $\partial_T w_\theta$  is computed via automatic differentiation with respect to  $T$ . Negative values indicate a violation and are penalized quadratically.

### 4.3.3 Butterfly-arbitrage penalty

Absence of butterfly arbitrage requires that the call price be convex in strike, i.e.,

$$\frac{\partial^2 C}{\partial K^2}(K, T) \geq 0,$$

which is equivalent to non-negativity of the risk-neutral density. I therefore define

$$L_{\text{butterfly}}(\theta) = \frac{1}{|\mathcal{G}_{\text{but}}|} \sum_{(K, T) \in \mathcal{G}_{\text{but}}} \left( \min \{0, \partial_K^2 C_\theta(S, K, T, r, q)\} \right)^2, \quad (6)$$

where  $\partial_K^2 C_\theta$  is obtained via second-order automatic differentiation of the Black–Scholes formula with respect to  $K$ . In practice, I evaluate this penalty on a dense grid of strikes and maturities that covers the support of the data.

## 4.4 Two-phase warm-up training strategy

Initial experiments revealed that if the physics penalties are activated from the first epoch with non-trivial weights, the optimizer often finds a nearly flat  $\sigma_\theta(K, T)$  that satisfies the derivative-based constraints at the expense of large pricing errors. To avoid this “mode collapse,” I adopt a two-phase schedule:

**Phase A (Warm-up).** For the first 30 epochs, I set  $\lambda_{\text{cal}} = \lambda_{\text{but}} = 0$  and train solely on the price loss  $L_{\text{price}}$ . This allows the network to learn the coarse shape of the implied-volatility smiles and term structures directly from data, unconstrained by the physics penalties.

**Phase B (Refinement).** Starting from epoch 31, I turn on the physics penalties with small weights:

$$\lambda_{\text{cal}} = \lambda_{\text{but}} = 10^{-4}.$$

Training continues up to epoch 150. In this phase, the optimizer nudges the data-fit solution towards one that also satisfies calendar and butterfly constraints, with only a modest increase in pricing error.

I train using mini-batch stochastic gradient descent with the Adam optimizer. Hyperparameters such as learning rate and batch size are tuned empirically to achieve stable convergence.

## 5 Experiments and Results

### 5.1 In-sample calibration on January 3, 2023

I first calibrate the model on SPY options from 3 January 2023 and evaluate performance on a held-out validation subset from the same day. Table 1 summarizes the main quantitative results.

Metric (Validation Set)	Phase A (Data Only)	Phase B (Final PINN)
Price Loss (MSE)	$\approx 2.09$	3.44
Constraint Loss	N/A (no constraints)	1.91
Pricing RMSE	—	\$1.76
MAE	—	\$1.23
Calendar-arbitrage violation	—	0.01% (on test grid)

Table 1: Performance before and after introducing physics-informed constraints.

Several observations emerge:

- The unconstrained Phase A model achieves a lower price MSE but exhibits numerous arbitrage violations, especially calendar violations at short maturities and convexity violations in the wings.
- After turning on the physics penalties, Phase B slightly increases the price MSE but drives the calendar-arbitrage violation rate on the evaluation grid down to zero and significantly reduces butterfly violations.
- The resulting pricing accuracy (RMSE around \$1.76, MAE around \$1.23) is acceptable for practical use, especially considering that many outliers are in illiquid, far-from-the-money contracts.

Qualitatively, the fitted IV smiles at fixed maturities are smooth and convex, passing through the center of the noisy market scatter cloud, with no visible oscillations. The 3D surface  $\sigma_\theta(K, T)$  appears smooth across both strike and maturity, without spikes or ridges.

### 5.2 Zero-shot generalization to January 4–5, 2023

To test whether the model learned structural features of the volatility surface rather than memorizing the Jan 3 cross-section, I evaluate the trained Phase B model directly on SPY option data from 4 and 5 January 2023, without any retraining. Using the same preprocessing pipeline, I compute the pricing RMSE on these new days:

- **Jan 4, 2023:** RMSE  $\approx 1.50$ ,
- **Jan 5, 2023:** RMSE  $\approx 0.93$ .

These zero-shot errors are comparable to or better than the in-sample RMSE, reflecting that the market conditions on these neighboring days are similar and that the model has captured the underlying term-structure and moneyness dependence robustly.

### 5.3 Transfer learning via short fine-tuning

In a production setting, one would recalibrate the model daily. To test whether the learned IV surface can be efficiently updated, I use the Jan 3 Phase B parameters as initialization and fine-tune on Jan 4 and Jan 5 data for only 10 additional epochs.

The resulting RMSEs are:

- **Jan 4, 2023:** RMSE  $\approx 1.47$  (about 2% improvement over zero-shot),
- **Jan 5, 2023:** RMSE  $\approx 0.87$  (about 7% improvement).

This shows that the model can be quickly adapted to new days with minor additional training, supporting an efficient daily recalibration workflow. The warm-started optimization converges rapidly and preserves the no-arbitrage properties inherited from the base model.

#### 5.4 Ablation: training without warm-up

For comparison, I train variants where  $\lambda_{\text{cal}}$  and  $\lambda_{\text{but}}$  are turned on from the beginning without a warm-up phase. In these runs, the optimizer often converges to:

- A nearly flat volatility surface with low variance across  $K$  and  $T$ ,
- Very small physics penalties (because derivatives vanish),
- But substantially higher pricing errors and poor fit to the observed smiles.

This confirms that the warm-up strategy is not just a cosmetic choice; it is critical to avoid trivial minima induced by the physics penalties and to achieve a good joint compromise between data fit and arbitrage-free structure.

## 6 Discussion and Conclusions

This project demonstrates that Physics-Informed Neural Networks provide a practical and effective framework for IV-surface inversion when combined with a suitable training strategy.

### 6.1 Key findings

The main findings are:

1. **Feasibility of PINN-based IV inversion.** A shallow MLP parameterizing  $\sigma(K, T)$ , coupled with the Black–Scholes formula as a differentiable layer, can recover a smooth IV surface with good pricing accuracy from realistic SPY option data.
2. **Importance of no-arbitrage penalties.** Without physics terms, the network fits noisy quotes but produces economically invalid surfaces with calendar and butterfly arbitrage. Introducing calendar and convexity penalties eliminates these violations on a dense grid while preserving most of the pricing accuracy.
3. **Warm-up training is crucial.** If constraints are active from epoch 0, the optimizer tends to settle on a trivial flat solution. A two-phase schedule that initially fits data and then gradually enforces constraints yields significantly better solutions.
4. **Robust temporal generalization.** The calibrated model generalizes well to neighboring trading days, both in zero-shot mode and under short fine-tuning, indicating that it has learned structural aspects of the SPY volatility surface.

### 6.2 Limitations and future work

Several limitations suggest natural extensions:

- **Scope of constraints.** I focus on static no-arbitrage constraints (calendar and butterfly). A more ambitious PINN would directly enforce the Dupire or Black–Scholes PDE as a residual term, possibly via an additional price network, at the cost of higher complexity.

- **Model class.** The current MLP may not fully capture extreme-wing behavior or very long maturities. Augmenting the architecture with explicit asymptotic priors or hybrid parametric–neural components could improve extrapolation.
- **Uncertainty quantification.** The inverse problem is not well-posed; multiple IV surfaces may fit the data similarly. Bayesian or ensemble approaches could provide uncertainty bands for  $\sigma(K, T)$ , highlighting regions where data are weak.
- **Multi-asset and intraday data.** Extending the approach to multiple underlyings or intraday time slices would test scalability and robustness of the framework in more complex settings.

Despite these limitations, the results suggest that PINNs are a promising tool for robust, arbitrage-aware volatility-surface calibration. The combination of neural-network flexibility, automatic differentiation, and financially motivated constraints yields surfaces that are both accurate and economically interpretable, and supports efficient re-calibration as market conditions evolve.

## Acknowledgments

This project was carried out as part of the AMCS 6045 course on inverse problems. I thank the instructor and classmates for discussions on inverse problems, regularization and PINNs, which influenced the design of this work.

## References

- Y. Aït-Sahalia and A. W. Lo. Nonparametric estimation of state-price densities implied by option prices. *Journal of Finance*, 53(2):499–547, 1998.
- R. Cont and P. Tankov. *Financial Modelling with Jump Processes*. Chapman and Hall/CRC, 2004.
- B. Dupire. Pricing with a smile. *Risk*, 7(1):18–20, 1994.
- J. Gatheral and A. Jacquier. Arbitrage-free SVI volatility surfaces. *Quantitative Finance*, 14(1):59–71, 2014.
- A. Hernandez. Model-free option pricing with deep learning. Working paper, 2017.
- M. Laurini. Imposing no-arbitrage conditions in implied volatility surfaces using constrained smoothing splines. Working paper, 2010.
- M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- X. Zhang, Y. Li, and Z. Wang. Physics-informed neural network for option pricing. arXiv preprint, arXiv:2312.06711, 2023.